

SPROCKET: Extending ROCKET to Distance-Based Time-Series Transformations With Prototypes

Nicholas Harner
Independent Researcher
Beavercreek, OH
memorialsignin@protonmail.com

December 8, 2025

Abstract

Classical Time Series Classification algorithms are dominated by feature engineering strategies. One of the most prominent of these transforms is ROCKET, which achieves strong performance through random kernel features. We introduce SPROCKET (**S**elected **P**rototype **R**andom **C**onvolutional **K**ernel **T**ransform), which implements a new feature engineering strategy based on prototypes. On a majority of the UCR and UEA Time Series Classification archives, SPROCKET achieves performance comparable to existing convolutional algorithms and the new MR-HY-SP (MultiROCKET-HYDRA-SPROCKET) ensemble's average accuracy ranking exceeds HYDRA-MR, the previous best convolutional ensemble's performance. These experimental results demonstrate that prototype-based feature transformation can enhance both accuracy and robustness in time series classification.

1 Introduction

Time series classification algorithms have undergone rapid development recently, including the introduction of a new family of algorithms based on convolutional kernels for feature extraction. ROCKET[6] exemplified these methods and achieved extremely high accuracy with minimal computation time. Rocket was refined into MiniROCKET[7] and MultiROCKET[19], which constrained the random selection of convolutional kernels and expanded the pooling operations, respectively. Convolutional methods were then expanded to include HYDRA[8], which incorporated dictionary-like methods for extracting time-series subsequences into a ROCKET-like classifier through competing convolutional kernels. The 2023 Bake-Off [16] established that an ensemble of MultiROCKET and Hydra was a top performing classifier, surpassed only by the HIVE-COTE

ensemble which requires substantially greater computational resources. However, these rankings may have changed with the development of ConvTran[10]. Convolutional feature extraction strategies have been remarkably successful in the time series domain. Their relative simplicity and applicability to both classical and deep learning methods make them a promising avenue for continued research.

Motivated by the success of combining dictionary-like approaches into convolutional methods, SPROCKET incorporates distance measures into ROCKET with minimal changes to the ROCKET architecture. Distances, especially elastic distances, have an extensive history in time series problems and are one of the algorithm families in the Bake-Off[16]. However, they are not trivially compatible with convolutional methods, since they require pairwise comparisons. SPROCKET addresses this issue through a randomized prototyping strategy to achieve baseline accuracy competitive with both ROCKET and HYDRA. Furthermore, the addition of SPROCKET to the HYDRA-MR ensemble improves its performance on a majority of the UCR and UEA Time Series Classification datasets [4] [5].

2 Intuition

Distances have been central to time series classification since the development of DTW and NN-DTW served as the baseline for the 2016 Bake-Off[3]. Distance-based approaches remain prominent in clustering [17] and time series specific distance measures continue to develop, with the recent introduction of amercing [12] and its incorporation into the Proximity Forest 2.0 [11]. As a robust nonparametric family of classification algorithms, distances have many attractive features. However, they face several fundamental limitations.

- Elastic distances are computationally expensive.

Elastic distance measures require a solution to a dynamic programming problem with computational complexity of $O(l^2)$, where l is the length of the series. The search space can be constrained to lower the complexity to $O(lw)$, where $w \ll l$.

- Distances capture only the some of a datasets properties.

Each distance imposes a single set of relations on a dataset. This single relation inherently limits the diversity of the learnable information from a dataset-distance pair.

The first issue can be addressed by selecting prototypes. Prototype based learning relies on representative samples of the dataset, called prototypes, to serve as references for all other points. For each prototype p and instance a , the distance $d(p, a)$ is a similarity measure between the pair. With a well selected prototype set, this signal can be used to derive useful information for all points in the dataset. Prototype selection reduces the computational burden of metric learning by limiting the number of points of comparison and their resultant distance calculations.

The issue of a limited derivable information given a dataset and a distance can be partially addressed by utilizing many different similarity measures and transformations of a time series. Taking derivatives is a classic example of this, as seen in DDTW, which applies derivatives to a series before taking the DTW distance. This transformation has been successful in the classification domain and the first derivative is incorporated into the Proximity Forest 2.0 as a baseline strategy for increasing ensemble diversity [11]. However the exceptional performance of convolutional kernels for time series classification suggests an additional approach. Apply the kernels from ROCKET as a feature transformation, then use a distance based approach in this new space. This observation motivates SPROCKET.

3 ROCKET Architecture and Modification

Starting with ROCKET is a reasonable first step in combining ROCKET and distances. The transform has one hyperparameter in its base form and only requires a specified number of kernels.

Algorithm 1 ROCKET

- 1: Set the number of kernels K for a training set X .
 - 2: **for** kernel $k_i = 1, 2, \dots, K$ **do**
 - 3: Parameterize kernel k_i , selecting a random set of weights, number of weights, and dilation.
 - 4: **for** training point p_j $j = 1, 2, \dots, |X|$ **do**
 - 5: Calculate $k_i(p_j)$.
 - 6: Calculate M pooling features (PPV, Max, and so on) for $k_i(p_j)$.
 - 7: **end for**
 - 8: **end for**
 - 9: Create matrix A of shape $(|X|, KM)$ and populate it with the pooling features.
 - 10: Fit a Cross Validated Ridge predictor on A .
-

The key steps of this algorithm are (i) applying kernels, (ii) pooling features, and (iii) using a model. Of these three, the pooling step is the natural place to modify the framework. This is also where prototype selection is required. Note that prototypes are essential here, since the predictor trained on the pooling features requires those features to be consistent. Even if $d(a, b) = d(c, d)$ these values could provide minimal information for a model depending on the relative positions of the points a, b, c, d in the space. In contrast, $d(a, b) = d(a, c)$ is far more likely to be meaningful, since the common point a is included in both features. The shared reference point should greatly increase the likelihood that features are informative for a model.

Now there are three remaining key design decisions. SPROCKET requires a distance measure, a prototype selection method, and a number of prototypes.

3.1 Number of Prototypes

Algorithm 2 SPROCKET

- 1: Set the number of kernels K and number of prototypes M for a training set X .
 - 2: **for** kernel $k_i = 1, 2, \dots, K$ **do**
 - 3: Parameterize kernel k_i as in ROCKET.
 - 4: **for** training point $p_j j = 1, 2, \dots, |X|$ **do**
 - 5: Calculate $k_i(p_j)$.
 - 6: **end for**
 - 7: Create a prototype set C_{k_i} for each kernel k_i , where $C_{k_i} \subset k_i(X)$, $|C_{k_i}| = M$.
 - 8: **for** training point $p_j j = 1, 2, \dots, |X|$ **do**
 - 9: Calculate the distances $d(c, k_i(p_j))$ between all prototypes in $c \in C_{k_i}$ and p_j .
 - 10: **end for**
 - 11: **end for**
 - 12: Create matrix A of shape $(|X|, KM)$ and populate it with the calculated distances.
 - 13: Fit a Cross Validated Ridge predictor on A .
-

3.1 Number of Prototypes

SPROCKET needs a number of prototypes, but without the other portions of the algorithm defined, the best heuristic is not obvious. However, we can describe two desirable properties of any method:

- The number of prototypes should be small.

ROCKET uses thousands of convolutional kernels in its default configuration. If the number of prototypes is large, SPROCKET will quickly become computationally infeasible as the number of distances calculations increases.

- The number of prototypes should increase with dataset size.

NN and other distance based learning methods are consistent, even if their convergence is slow. As datasets increase in size, they could exhibit richer structure that requires more prototypes to capture.

These properties compete. As a compromise, the number of prototypes can be set to a monotonic, but slowly increasing function. Any logarithmic function on the size of the dataset would achieve this goal, but for experimental stability it should be consistent. In this paper, the size of the prototype set is always

$$\lceil \log_4(|X|) \rceil \quad (1)$$

rounding up to ensure a minimum number of prototypes for small datasets. The datasets tested in this paper range from 10 – 5000 training points, which results in 2 – 7 prototypes per kernel. Alternative strategies are possible and should be explored, but this paper will only demonstrate a minimal baseline algorithm.

3.2 Prototype Selection

Prototypes must be chosen, but most principled forms of prototype selection require computing many pairwise distances. Since the number of computed distance measures dominate the runtime differential between SPROCKET and other convolutional algorithms, any increases to computational load should be carefully considered.

Two baseline prototype selection algorithms require no additional distance calculations.

1. Uniform Random Sampling of prototypes from the training set.
2. Stratified Random Sampling of prototypes from the training set in proportion to their class prevalence.

Both are inexpensive, unbiased, scalable, simple to implement, and standard solutions to this issue.

A natural alternative approach is to perform all or part of a clustering algorithm that provides centroids or medoids. For instance, K-Means++ [1] has natural applications to time series clustering and efficient variants such as KASBA [13] for time series specific measures. The initialization step of Kmeans++ is:

Algorithm 3 Kmeans++ Initialization

- 1: Add one center c to the center set C , chosen uniformly at random among the data points X .
 - 2: For each data point $x \in X \setminus C$, compute $D(x, C)$, or the distance between x and the nearest center in C .
 - 3: Add c_{new} to C , where c_{new} is chosen with probability proportional to $D(x, C)$ for all $x \in X \setminus C$.
 - 4: Repeat Steps 2 and 3 until k centers have been chosen.
-

Initialization therefore requires, at minimum $|X| - 1$ distance calculations. In the worst case, each center added to C will require $|X| - |C|$ distance calculations. Exploiting the triangle inequality can avoid distance calculations if the distance function is a metric, but the worst case initialization of KMeans++ will approximately double the required number of distance calculations relative to random selection. This does not include calculations for subsequent iterations Kmeans itself, which must in turn be applied for all of the hundreds or thousands of kernels in a ROCKET framework.

The goal of this study is to establish a minimum baseline to combine distance learning with ROCKET. Although prototype selection offers potential to improve accuracy, it comes with substantial computational cost above random selection. Therefore, this paper will confine prototype selection strategies to random sampling, leaving exploration of alternatives to future work.

3.3 Distance Measures

Any design requires a set of distance measures used for prototype features. There are many possible measures to consider, but the distance measures present in the Aeon toolkit [15] form a readily implementable baseline for this study. For detailed descriptions, we refer the reader to [14]. The selected measures for this study are:

Long Name	Abbreviation	Notes
Euclidean	Euclidean	Baseline non-elastic measure
Dynamic Time Warping	DTW	Baseline elastic distance
Weighted DTW	WDTW	DTW with a weighted warping penalty
Amerced DTW	ADTW	DTW with amercing
Edit Distance with Real Penalty	ERP	Levenshtein-style distance with real valued data
Time Warp Edit	TWE	Edit distance with explicit time-warping penalty
Move-Split-Merge	MSM	Edit-based with local warping penalties

Table 1: Elastic and non-elastic distance measures evaluated.

This set is not exhaustive but does provide a diverse baseline. Exploring additional measures will be left to later work.

4 Algorithm Analysis

We will now analyze SPROCKET’s computational complexity to determine its scaling factors. Because pairwise distance calculations dominate the runtime of the transform, this analysis will focus on those.

The SPROCKET transform creates $\lceil \log_b(|X|) \rceil$ prototypes for each kernel, where b is a constant (set to 4 for our experiments), and k kernels. Now let the dataset have $|X|$ instances, m channels, and length l on each channel. For multivariate series, let the distances be calculated independently on each channel. Then, the number of distance calculations N_d to perform the transform is:

$$N_d = k \cdot n \cdot \log_b(n) \cdot m \quad (2)$$

If an elastic distance measure with a Sakoe-Chiba band $w \in 2, 3 \dots l$ is used, then a single distance measure costs $O(lw)$. So SPROCKET transform, excluding

the convolutional kernel computational cost inherited from ROCKET costs

$$O(k \cdot n \cdot \log_b(n) \cdot m \cdot D) = O(k \cdot n \cdot \log_b(n) \cdot m \cdot l \cdot w) \quad (3)$$

For non-elastic distances such as the Euclidean distance, the runtime is reduced by a factor of w , since the distance calculation has $O(l)$. This complexity equation can then be compared to the Proximity Forest’s time complexity, which is:

$$O(n \cdot \log_2(n) \cdot m \cdot l \cdot w \cdot r \cdot c \cdot t) \quad (4)$$

Where c is the number of classes, r the number of candidate splits per node, and t the number of trees. SPROCKET will have lower complexity in the cases where

$$r \cdot c \cdot t > k \wedge b > 2 \quad (5)$$

Since SPROCKET’s scaling is independent of the number of classes, it will hold a relative scaling advantage over other distance based methods in multi-class problems. It will also be at a relative advantage in cases where the Proximity Forest requires many splits per node, assuming the construction of the resulting classifier is computationally negligible relative to feature extraction.

5 Experimental Results

5.1 Settings

For all experiments, parameters were held constant unless otherwise noted. All elastic distances were calculated with a window parameter for the Sakoe-Chiba band set to $\lfloor \sqrt{l} \rfloor$, features were prepared in parallel with 20 threads, and elastic distance parameters were otherwise untuned, using the Aeon-Toolkit’s default settings.

The ROCKET variant classification models were trained with scikit-learn’s RidgeClassifierCV() classifier with the default settings.

All experiments were conducted on a System76 Adder WS running PopOS, with 16 GB of Memory, and 32 13th Gen Intel® Core™ i9-13900HX cores.

5.2 Prototype and Distance Selection

From Section 3, we have two unresolved design concerns:

- Which distances to use?
- How to select prototypes?

We can start with a dual mode experiment. We will evaluate the SPROCKET algorithm on a representative subset of the UCR archive twice for each distance measure. The first time, we will select the prototypes completely randomly and the second time we will use stratified random sampling. This will yield two tests of the relative rankings for each distance measure and seven tests of the

5.2 Prototype and Distance Selection

prototyping strategy, one for each distance measure. The selected datasets span a variety of applications and lengths, a full description is included in Appendix A. For this test, we set the number of kernels to be 512 for all distance measures.

First, we can consider the two relative results for distance measures with true and stratified randomly sampled prototypes.




Figure 1: Relative comparisons of Distance Measures with the SPROCKET Transform.




Figure 2: Stratified Results maintain a similar rank ordering to true random.

We can infer some information from these tests about the relative performance of the distance measures. The edit-based distances all perform well, with TWE, MSM, and ERP maintaining the first, second, and third best performing distances in each set of tests. Meanwhile, the warping based distances underperform, with WDTW and DTW failing to outperform the Euclidean baseline on average and ADTW outperforming Euclidean distances only once. However, all of the warping distances are the top performing distances on at least two datasets in both tests and all distances are the top performers on several datasets. We

5.2 Prototype and Distance Selection

can therefore conclude that the warping distances are not strictly dominated in accuracy, but do perform worse on average than the edit based distances.

We can also examine the relative time each distance takes to prepare the prototype features.




Figure 3: Total Feature Transformation Time For Each Distance, Random Prototypes.

Here, we have four groups. We have the baseline times for all the warping based distances and TWE. The other two edit based distances take longer, MSM taking 2.8 times longer than the baseline set of $\{EDR, DTW, ADTW, WDTW\}$ and TWE taking 4.5 times baseline. Finally, the Euclidean distance is two orders of magnitude faster than all elastic distances, taking under 0.01 of the base time. Unfortunately, on these datasets the most accurate distances are also the most computationally expensive by substantial factors.

Finally, the distances could be valuable in an ensemble. To examine this possibility, we can consider several ensemble statistics for each distance's predictions.

All of the distances are highly correlated with each other on these datasets, with a lowest correlation 0.70 (between DTW and WDTW with Stratified Random Prototypes) and a highest correlation of 0.89 (between MSM and TWE with Stratified Random Prototypes). The other ensembling statistics are similarly discouraging. We should also note the high correlation between the edit distances. They are far more correlated with each other than with the warping based distances, suggesting that the relative performance improvements from ensembling them will be limited on similar data, even though they are relatively more accurate than the warping based distances.

Finally, we can compare the two prototype selection strategies. The table

5.2 Prototype and Distance Selection




Figure 4: Ensemble Grid for Random Prototypes.

5.2 Prototype and Distance Selection

Stratified Random Prototypes Ensemble Analysis Grids




Figure 5: Ensemble Grid for Stratified Random Prototypes.

5.3 Ensemble Experiments

below counts the number of times Random Selection and Stratified Random Selection outperformed on the 33 datasets in our small sample for each distance measure.

Distance	Random Wins	Stratified Wins	Ties
Euclidean	17	12	4
DTW	14	15	4
ADTW	16	14	3
WDTW	12	14	7
ERP	17	12	4
MSM	14	14	5
TWE	18	11	4
Total	108	92	31

Table 2: Relative Comparisons of Random and Stratified Random Accuracies on 33 Datasets with 7 Distance Measures.

These results suggest a weak advantage for true random sampling over stratified random sampling on the UCR datasets. Performing a one-sided sign test to determine whether random sampling is preferable to stratified random sampling with the 200 total nontied trials results in a p-value of 0.14 in favor of true random. Though this is not statistically significant, combining these results with the relative simplicity advantage of true random sampling over stratified suggests that true random is acceptable for this initial exploratory analysis of SPROCKET. We used true random sampling as the default for the remainder of this paper.

5.3 Ensemble Experiments

The previous section suggested that combining multiple distances is unlikely to increase SPROCKET accuracy or reduce computation time for SPROCKET due to the high correlation of distance measures. However, ensembling is well established in the time series distance learning literature and we wished to test whether it would help even in this case. To test this possibility, we conducted a limited test of potential ensembles.

To begin, we divided the distance measures into three groups based on their conceptual framework and the empirical results of our experiments.

Then we adopted an ensembling strategy. Ideally, we would validate the various distance measures on holdout data and form an ensembling strategy based on those results, as seen in the Elastic Ensemble[2]. But that is not necessary for the other convolutional algorithms and would incur substantial additional computational requirements. Instead, for a simple, default version of SPROCKET, we will use a constant fraction of the allotted convolutional kernels and see which ensemble, if any, performs the best.

We tested six ensembles and seven unensembled distance measures on the same subset of the UCR Benchmark in the previous sections. All classifiers will

5.3 Ensemble Experiments

Label	Distances	Accuracy	Time	Correlation
Inelastic	Euclidean	Moderate	Very Fast	High
Warping	ADTW, DTW, WDTW	Moderate	Slow	High
Edit	TWE, MSM, ERP	High	Slow–Very Slow	Very High

Table 3: Distance sets and kernel characteristics.

have either 600 or 1200 kernels.

Label	Distances	Kernel Counts
Top2	TWE, ADTW	300, 300
Top2+e	TWE, ADTW, Euclid	300, 300, 600
Top4	TWE, ADTW, MSM, DTW	150, 150, 150, 150
Top4+e	TWE, ADTW, MSM, DTW, Euclid	150, 150, 150, 150, 600
All Elastic	TWE, ADTW, MSM, DTW, ERP, WDTW	100, 100, 100, 100, 100, 100
All	TWE, ADTW, MSM, DTW, ERP, WDTW, Euclid	100, 100, 100, 100, 100, 100, 600

Table 4: Distance sets and kernel counts.

These ensemble combinations of distances were chosen to consider contributions from each of the three groups, balance computational considerations, and manage experimental complexity. We will consider equal contributions from each of the two Elastic distance groups, and including the most accurate, two most accurate, and all distances from each of those groups. The three best performing distances in the previous section were all highly correlated and had Q-statistics between 0.86 – .089. This suggests that within group ensembling will be of limited utility and, given the computational limitations at hand, are not worth exploring in this initial assessment. Instead, this test focuses on cross group diversity. There are two equal contributions from each of the two Elastic distance groups, taking the two most accurate from each elastic group in Top2, the four most accurate in Top4, and all distances from each of those groups. Finally, since the Euclidean distance requires orders of magnitude less computation, it can be included with minimal additional computational cost on top of the Elastic distance kernels.

The top performing ensembles on this test (Top2 and Top2+e) were both outperformed by all edit based versions of SPROCKET (TWE, MSM, and ERP). Top2+e held the same average rank as the Euclidean SPROCKET, and Top2 held a slightly worse average rank, though the difference was small at 0.06. Both

5.3 Ensemble Experiments




Figure 6: Relative Performance of Ensembles and Base Classifiers.

also achieved the best possible rank more times than the Euclidean SPROCKET, with 3 and 4 first ranks as compared to the Euclidean 2, but less than all edit distances, with 8, 12, and 19 best ranks for ERP, MSM, and TWE respectively.

We can also examine the ensembling statistics for all tested classifiers.

All Ensembles have higher correlations with each other than the single distance classifiers. The most correlated any ensemble is with any single classifier is 0.80, which is also the least correlated any ensemble is to any other ensemble. Simultaneously, Ensemble correlations can range as high as 0.90, which suggests that the ensembles are more similar to each other than any of their components.

The relative under performance of ensembling multiple distance measures with a naive strategy as opposed to simply using the best distances suggests that hyperparameter tuning may be required to achieve high ensemble performance with measures. The simple concatenated ensembles tested here did not gain significant advantage from diverse distance measures on this dataset. Since more complex ensembling such as stacking and weighted voting require additional computational work for validation, we will not consider it further in this paper.

Instead, we will adopt MSM as the default elastic distance metric for this paper. Although the TWE metric has outranked MSM or tied MSM in average rank on all tests, none of those have been statistically significant. However, the computational requirements for TWE are much larger than MSM, with the MSM transformation taking 63% on average, of TWE's transformation time.

Furthermore the large computational difference between the Euclidean distance and all other distances suggests that the Euclidean distance should still be considered separately for computationally constrained applications. It also performs relatively well, outranking WDTW in all tests and DTW in some. In light of this, we can consider two distances, one elastic and one inelastic, as default SPROCKET recommendations.

5.3 Ensemble Experiments




Figure 7: Ensemble Statistics for Tested Ensembles.

5.4 Kernel Scaling

5.4 Kernel Scaling

SPROCKET is a convolutional transform and other convolutional transforms are sensitive to the K hyperparameter that determines the number of kernels used. We therefore examine SPROCKET’s sensitivity to K in a limited test on the 33 dataset subset before performing a large-scale evaluation on the UCR/UEA benchmark.

To test this, we created MSM classifiers where K is set to every power of two between 8 and 4096 and tested them on the small subset of UCR that we previously used.




Figure 8: Relative Performance of Power of Two Kernel Counts for SPROCKET. Note that $\text{msm}8$ has 8 kernels, $\text{msm}16$ has 16, etc.

The Critical Difference diagram reveals that the number of kernels does improve average performance, but not fully uniformly. $\text{MSM}1024$ has a higher average rank than $\text{MSM}512$ by .11, a small but notable result. The higher kernel counts are also more frequently the best performing classifiers, but again this is not uniform.

The time statistics are far more uniform and show a clear linear scaling pattern for various values of K .

Finally, correlations and Q-Statistics between different kernel counts of MSM-SPROCKET increase for larger values of K , while disagreements and double faults decrease. These patterns in ensemble statistics suggest diminishing marginal diversity among the independently parameterized MSM-SPROCKET variants. Such diminishing diversity reflects relatively muted gains and expected ensembling accuracy improvements from increased kernel counts as K increases.

$\text{MSM}512$ is the smallest version of SPROCKET where Q-Statistics plateau at 0.90 for larger versions of SPROCKET. This suggests that $\text{MSM}512$ is at the point of diminishing returns for larger K . It also has an identical kernel count to HYDRA, another prominent transform that hybridizes convolutional and non-convolutional approaches. This makes $K = 512$ a pragmatically attractive default setting for K . This small scale testing on 33 datasets cannot definitively

5.5 Large Scale MSM Benchmarking




Figure 9: Total time to transform data for various MSM kernel counts.

establish the optimal K threshold and that threshold could be dataset specific, but 512 kernels balance accuracy and computational requirements within the tested range.

These results indicate that while increasing K generally improves performance, the gains are not linear with increasing values of K, while corresponding increases in computation time are linear. For large-scale evaluation, we therefore adopt K=512 as a balanced default.

5.5 Large Scale MSM Benchmarking

We now proceed to a large-scale test of SPROCKET on the UCR Time Series Classification Library. We tested MultiROCKET, HYDRA, SPROCKET, MR-HY (MultiROCKET-HYDRA), MR-SP (MultiROCKET-SPROCKET), HY-SP (HYDRA-SPROCKET), MR-HY-SP (MultiROCKET-HYDRA-SPROCKET), and QUANT[9], as a non-convolutional baseline. The MultiROCKET, HYDRA, and SPROCKET ensembles were formed by concatenating their transformed features and using a single RidgeCV classifier, as is standard for the existing MR-HYDRA ensemble. All non-SPROCKET transforms utilized their standard configurations in the Aeon-Toolkit and the tested SPROCKET configuration utilizes MSM distances, 512 kernels, $\lceil \log_4(|X|) \rceil$ random prototypes, and a window parameter of $\lfloor \sqrt{l} \rfloor$, where l is the length of the series.

For this test, we sought to capture a wide selection of datasets from the UCR Classification Benchmark Library. We excluded:

- Datasets with unequal length series, which are not compatible with any of the tested classifiers.

5.5 Large Scale MSM Benchmarking




Figure 10: Ensemble Statistics for Powers of Two Scaled MSM Sprockets.

5.5 Large Scale MSM Benchmarking

- Datasets with more than 5000 training instances. Due to the large number of features produced by MultiROCKET, those datasets would require stochastic gradient descent (SGD) on available hardware instead of closed-form Ridge regression used elsewhere, reducing comparability.
- Datasets with length greater than 500. Since SPROCKET scales superlinearly with length, a maximum length was needed.
- Datasets with length under 9. Since MultiROCKET requires datasets to have a length of at least 9, this is to maintain comparability between all classifiers.

This resulted in 98 datasets, which are detailed in Appendix A. These datasets were tested 5 times with different random seeds and the standard train-test split provided by the UCR datasets. The results were averaged and presented below.




Figure 11: Aggregate CD and Best Performances

These results establish that, although SPROCKET is the worst performing in individual algorithm, it ensembles at least as well as HYDRA and MultiROCKET for convolutional algorithms. It also achieves the best rank on 10 of the 98 datasets, showing that it is not strictly dominated as an individual algorithm. More importantly, the SPROCKET ensembles perform very well. The MR-HY-SP ensemble outranks all other algorithms in this test and the MR-SP ensemble outranks MR-HY, despite the fact that HYDRA outranks SPROCKET individually. The large number of individual best ranks for QUANT are likely attributable to it being the only nonlinear algorithm in this sample, but are still notable.

Detailed head-to-head comparison scatterplots of each algorithm are not included here because of their high number. To examine them in detail, please see Appendix B.

We can also examine the ensembling statistics for the tested algorithms.

The Pairwise Q-Statistics for MultiROCKET-SPROCKET, MultiROCKET-HYDRA, and HYDRA-SPROCKET are 0.64, 0.64, and 0.62 respectively. This

5.5 Large Scale MSM Benchmarking




Figure 12: Ensembling Statistics for Large Scale MSM Test

5.5 Large Scale MSM Benchmarking

relationship may explain why the pairwise ensembles all perform similarly. All the ensembles containing MultiROCKET are more correlated with MultiROCKET than any of their other subcomponents, suggesting that MultiROCKET is dominant subcomponent in the ensembles. Finally, QUANT retains similar ensembling statistics to each convolutional method, which suggests that the predictions produced by the covolutional methods are still meaningfully different than at least one non-convolutional method.

We can now turn to the running times. By taking the average transformation time of the SPROCKET transforms against the predicted value from our theoretical analysis in Section 4, we can confirm that the theoretical results are somewhat representative of our observed results, which is to be expected due variance caused by randomization in the SPROCKET algorithm and hardware implementation details.




Figure 13: SPROCKET Predictions Vs Observed Times

Finally, we can see that the SPROCKET transform with the MSM distance is substantially more computationally expensive than the other compared transforms.

5.6 Large Scale Euclidean Benchmarking

Algorithm	Total Time (Seconds)
SPROCKET	24353
MultiROCKET	364
HYDRA	193
QUANT	1373
MR-HY	590
MR-SP	24657
HY-SP	24415
MR-HY-SP	24796

Table 5: Total Time Taken For all Tests in Large Scale MSM Testing.

5.6 Large Scale Euclidean Benchmarking

Because of the large computational difference between MSM distance SPROCKET and the other convolutional algorithms, we repeated the large scale benchmark with the Euclidean distance measure. The experimental setup was otherwise identical. All non-SPROCKET transforms utilized their standard configurations in the Aeon-Toolkit and the tested SPROCKET configuration utilizes Euclidean distances, 512 kernels, $\lceil \log_4(|X|) \rceil$ random prototypes, and no window parameter, since it is not required for the Euclidean distance.

The experiment was repeated 5 times with different random seeds and the average results are presented below.




Figure 14: Euclidean Aggregate CD and Best Performances

Although the ranking is less clear cut than in the MSM case, a similar ranking is present in for the Euclidean SPROCKET. The MR-HY-SP ensemble holds the best possible rank, though the difference between it and the other algorithms has decreased. Interestingly, the MR-HY-SP algorithm has more overall best finishes in this configuration, even though the sum of its ranks has increased. Furthermore, the HY-SP ensemble is now ranked under QUANT, but the gap is

5.6 Large Scale Euclidean Benchmarking

not large.

Detailed head-to-head comparison scatterplots of each algorithm are again in Appendix B.

Aggregate Euclidean SPROCKET Ensemble Analysis Grids




Figure 15: Ensembling Statistics for Large Scale Euclidean Test

The Pairwise Q-Statistics for MultiROCKET-SPROCKET, MultiROCKET-HYDRA, have decreased to 0.59 and 0.55, which may explain why the accuracy decrease is relatively small under ensembling. Euclidean SPROCKET is slightly more independent than MSM SPROCKET from existing convolutional algorithms, though it is less accurate.

We can also repeat our test of predicted vs observed running times. Note that our prediction function has changed, since the Euclidean distance contains no windowing parameter and is instead strictly linear on length.

Finally, we can see that the Euclidean SPROCKET transform is much more computationally efficient than the MSM version, with the lowest total computation time of any tested algorithms.

5.6 Large Scale Euclidean Benchmarking




Figure 16: Euclidean SPROCKET Predictions Vs Observed Times

Algorithm	Total Time (Seconds)
SPROCKET	169
MultiROCKET	437
HYDRA	212
QUANT	1325
MR-HY	778
MR-SP	668
HY-SP	254
MR-HY-SP	819

Table 6: Total Time Taken For all Tests in Large Scale Euclidean Testing.

6 Future Work

This paper introduces SPROCKET as a proof of concept for prototype driven integration of distance based time series methods into ROCKET. This initial step leaves several promising avenues for continued research and optimization.

6.1 Improvements on SPROCKET Design

The design process employed in this research is limited. It did not examine many design decisions, including but not limited to:

- Non-random prototype selection.
- Weighted or adaptive distance measure ensembles.
- The $\lceil \log_4(|X|) \rceil$ prototype scaling heuristic.
- Pre and post convolution transformations (normalization, differencing, scaling, etc).
- Other distance measures beyond the seven evaluated.

All of these could be potentially improved and represent avenues for further research.

6.2 Improvements from Convolutional Transforms

This SPROCKET baseline uses the same convolutional architecture as baseline ROCKET, and does not include subsequent refinements since ROCKET’s introduction. Convolutional strategy developments may be of interest and include:

- First order difference transforms before kernel transformations, as seen in the MultiROCKET.
- Constrained kernel randomization to improve computation, feature extraction efficiency, or consistency as seen in limited kernel randomization in Mini and MultiROCKET.

6.3 Improvements from Distance Classifiers

- Pruning the kernels after initial training. Several strategies exist to perform this in the ROCKET framework, including Detach-ROCKET [20] and S-ROCKET [18]. The computational improvements from reducing distance calculations will be higher for SPROCKET than for other ROCKET variants, due to the higher cost of distance calculations per kernel.

6.3 Improvements from Distance Classifiers

SPROCKET is a hybrid Convolutional-Distance transformation and may benefit from the strategies employed by distance based time series classifiers. These include:

- Derivative transformations before or after kernel transformations, as seen in the Proximity Forest 2.0. This transformation is similar to the transformation use in MultiROCKET.
- Combining HYDRA and Minkowski distances as a non-elastic distance measure, as in the Proximity Forest 2.0
- Early abandonment and other computational optimizations of distance calculations to reduce runtime without sacrificing accuracy.

Any of these could improve both the accuracy and computational efficiency of SPROCKET.

6.4 Comprehensive Empirical Characterization

The empirical evaluation in this paper was constrained by computational resources and focused on datasets with moderate length (≤ 500) and size (≤ 5000 instances). Several important questions remain:

- Performance on long time series (length > 500), potentially using approximation methods or sampling strategies for distance computation.
- Scalability to large datasets ($> 5,000$ instances) and testing on the Monash MONSTER benchmark, utilizing SGD or other scaling methods.
- Sensitivity analysis for key hyperparameters beyond kernel count K , including prototype count, window parameters, and distance-specific settings.
- Evaluation on extrinsic regression and forecasting tasks to extend beyond simple classification to broader time series learning problems.
- Performance on domain-specific benchmarks (e.g., audio classification, medical time series, financial datasets).
- Stability and variance across different random initializations, particularly for datasets where SPROCKET shows high variability.
- Theoretical analysis of SPROCKET’s approximation properties and relationship to nearest-neighbor classification and prototype methods.

6.5 Implementation and Accessibility

The reference implementation of SPROCKET was adapted from existing transformations in Aeon and inherits SciKit-Learn compatibility. It is available at https://github.com/username76543/sprocket_public. This base level of integration could be built on in several ways:

- Optimized implementations leveraging additional parallelization and vectorization, beyond the numba optimizations included in the reference code.
- GPU acceleration for distance computations, particularly for large-scale applications.
- Integration into popular time series classification libraries such as aeon, sktime, and tslearn. The provided reference code was built utilizing aeon standards, but support should be provided for other popular libraries.
- Automated hyperparameter selection tools to reduce manual tuning burden.
- Comprehensive documentation and usage examples for different application domains.

While SPROCKET is not individually the top performing convolutional transform, its contribution to convolutional ensemble learning is demonstrably valuable and represents an improvement on existing methods. Further developments could build on this contribution to offer new fast and accurate convolutional methods. We invite the time series research community to explore any of these proposed research directions or others not listed.

7 Conclusion

This paper demonstrates that distance-based prototype features can be effectively integrated into convolutional time series classification frameworks. SPROCKET, while not individually state-of-the-art, substantially improves ensemble performance: the MR-HY-SP ensemble achieves the best average rank across 98 UCR benchmark datasets, surpassing previous convolutional ensembles. This improvement validates that complementary predictions from different algorithmic families enhance classification accuracy beyond what individual methods achieve.

The computational cost of elastic distance calculations presents a tradeoff: MSM-based SPROCKET requires significantly more computation than other convolutional methods and Euclidean SPROCKET reduces this by two orders of magnitude. Both significantly improve convolutional ensembles, but selecting the appropriate measure will depend on application constraints. The simple baseline design employed here—random prototypes, logarithmic scaling prototype numbers, standard distances—leaves substantial room for optimization through improved prototype selection, kernel pruning, and computational optimization.

By bridging convolutional and distance-based approaches, SPROCKET opens a productive research direction at the intersection of two major time series classification paradigms.

8 Citations

References

- [1] David Arthur and Sergei Vassilvitskii. k-means++: The advantages of careful seeding, 2007.
- [2] Anthony Bagnall. Time series classification with ensembles of elastic distance measures. *Data Mining and Knowledge Discovery*, 29, 06 2014.
- [3] Anthony Bagnall, Aaron Bostrom, James Large, and Jason Lines. The great time series classification bake off: An experimental evaluation of recently proposed algorithms. extended version, 2016.
- [4] Anthony Bagnall, Hoang Anh Dau, Jason Lines, Michael Flynn, James Large, Aaron Bostrom, Paul Southam, and Eamonn Keogh. The uea multivariate time series classification archive, 2018, 2018.
- [5] Hoang Anh Dau, Anthony Bagnall, Kaveh Kamgar, Chin-Chia Michael Yeh, Yan Zhu, Shaghayegh Gharghabi, Chotirat Ann Ratanamahatana, and Eamonn Keogh. The ucr time series archive, 2019.
- [6] Angus Dempster, François Petitjean, and Geoffrey I. Webb. Rocket: exceptionally fast and accurate time series classification using random convolutional kernels. *Data Mining and Knowledge Discovery*, 34(5):1454–1495, July 2020.
- [7] Angus Dempster, Daniel F. Schmidt, and Geoffrey I. Webb. Minirocket: A very fast (almost) deterministic transform for time series classification. In *Proceedings of the 27th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, KDD ’21, page 248–257. ACM, August 2021.
- [8] Angus Dempster, Daniel F. Schmidt, and Geoffrey I. Webb. Hydra: Competing convolutional kernels for fast and accurate time series classification, 2022.
- [9] Angus Dempster, Daniel F. Schmidt, and Geoffrey I. Webb. Quant: A minimalist interval method for time series classification, 2023.
- [10] Navid Mohammadi Foumani, Chang Wei Tan, Geoffrey I. Webb, and Mahsa Salehi. Improving position encoding of transformers for multivariate time series classification. *Data Mining and Knowledge Discovery*, 38(1):22–48, September 2023.

-
- [11] Matthieu Herrmann, Chang Wei Tan, Mahsa Salehi, and Geoffrey I. Webb. Proximity forest 2.0: A new effective and scalable similarity-based classifier for time series, 2023.
 - [12] Matthieu Herrmann and Geoffrey I. Webb. Amercing: An intuitive, elegant and effective constraint for dynamic time warping, 2021.
 - [13] Christopher Holder and Anthony Bagnall. Rock the kasba: Blazingly fast and accurate time series clustering, 2024.
 - [14] Christopher Holder, Matthew Middlehurst, and Anthony Bagnall. A review and evaluation of elastic distance functions for time series clustering. *Knowledge and Information Systems*, 66(2):765–809, September 2023.
 - [15] Matthew Middlehurst, Ali Ismail-Fawaz, Antoine Guillaume, Christopher Holder, David Guijo Rubio, Guzal Bulatova, Leonidas Tsaprounis, Lukasz Mentel, Martin Walter, Patrick Schäfer, and Anthony Bagnall. aeon: a python toolkit for learning from time series, 2024.
 - [16] Matthew Middlehurst, Patrick Schäfer, and Anthony Bagnall. Bake off redux: a review and experimental evaluation of recent time series classification algorithms. *Data Mining and Knowledge Discovery*, 38(4):1958–2031, April 2024.
 - [17] John Paparrizos, Fan Yang, and Haojun Li. Bridging the gap: A decade review of time-series clustering methods, 2024.
 - [18] Hojjat Salehinejad, Yang Wang, Yuanhao Yu, Tang Jin, and Shahrokh Valaee. S-rocket: Selective random convolution kernels for time series classification, 2022.
 - [19] Chang Wei Tan, Angus Dempster, Christoph Bergmeir, and Geoffrey I. Webb. Multirocket: Multiple pooling operators and transformations for fast and effective time series classification, 2022.
 - [20] Gonzalo Uribarri, Federico Barone, Alessio Ansuini, and Erik Fransén. Detach-rocket: Sequential feature selection for time series classification with random convolutional kernels, 2024.

Appendix A Datasets

Two sets of datasets were used for the empirical evaluation. All were taken from the UCR Time Series Classification Archive and the UEA Multivariate Time Series Archive. For all tests, they were imported in their standard Train/Test splits via the Aeon Toolkit.

The first, smaller set of datasets were used in the parameter evaluation studies. These were selected to cover a variety of train and test sizes, dimensions, number

of classes, and application areas, while allowing for faster iteration through a reduced sampling space.

Dataset Name	Train Size	Test Size	Length	Classes
Coffee	28	28	286	2
Beef	30	30	470	5
FaceFour	24	88	350	4
Lightning7	70	73	319	7
Wine	57	54	234	2
Meat	60	60	448	3
Ham	109	105	431	2
ECG200	100	100	96	2
ItalyPowerDemand	67	1029	24	2
MoteStrain	20	1252	84	2
SonyAIBORobotSurface1	20	601	70	2
SonyAIBORobotSurface2	27	953	65	2
TwoLeadECG	23	1139	82	2
GunPoint	50	150	150	2
Trace	100	100	275	4
CBF	30	900	128	3
SyntheticControl	300	300	60	60
DiatomSizeReduction	16	306	345	4
SwedishLeaf	500	625	128	15
Wafer	1000	6164	152	2
ECG5000	500	4500	140	5
FaceAll	560	1690	131	14
FacesUCR	200	2050	131	14
WordSynonyms	267	638	270	25
Adiac	390	391	176	37
ChlorineConcentration	467	3840	166	3
Yoga	300	3000	426	2
MedicalImages	381	760	99	10
FordA	3601	1320	500	2
FordB	3636	810	500	2
ElectricDevices	8926	7711	96	7
SpokenArabicDigits	6599	2199	93	10
Handwriting	150	850	152	26

Table 7: Datasets used for Parameter Selection.

The larger set of datasets was also selected from the UCR archive, but chosen more broadly to include 98 of the 132 datasets available. It includes all datasets included in the previous test except for Electric Devices and SpokenArabicDigits, which would have required a switch to SGD from standard Ridge Regression to

fit in available hardware memory. This switch may have affected comparability and necessitated excluding the datasets. The additional datasets included are listed below.

[H]

Dataset Name	Train Size	Test Size	Length	Classes
DuckDuckGeese	60	40	270	5
PEMS-SF	267	173	144	7
MindReading	727	653	200	5
PhonemeSpectra	3315	3353	217	39
ShapeletSim	20	180	500	2
EMOPain	1093	50	180	3
SmoothSubspace	150	150	15	3
MelbournePedestrian	1194	2439	24	10
Chinatown	20	345	24	2
JapaneseVowels	270	370	29	9
RacketSports	151	152	30	4
LSST	2459	2466	36	14
Libras	180	180	45	15
FingerMovements	316	100	50	2
NATOPS	180	180	51	6
SharePriceIncrease	965	965	60	2
ERing	30	270	65	6
PhalangesOutlinesCorrect	1800	858	80	2
ProximalPhalanxOutlineCorrect	600	291	80	2
MiddlePhalanxOutlineCorrect	600	291	80	2
DistalPhalanxOutlineCorrect	600	276	80	2
ProximalPhalanxTW	400	139	80	2
ProximalPhalanxOutlineAgeGroup	400	205	80	3
MiddlePhalanxOutlineAgeGroup	400	154	80	3
MiddlePhalanxTW	399	154	80	6
DistalPhalanxTW	400	139	80	6
DistalPhalanxOutlineAgeGroup	400	139	80	3
BasicMotions	40	40	100	4
TwoPatterns	1000	4000	128	4
BME	30	150	128	3
EyesOpenShut	56	42	128	2
ECGFiveDays	23	861	136	2
ArticularyWordRecognition	275	300	144	25
PowerCons	180	180	144	2
Plane	105	105	144	7
GunPointOldVersusYoung	135	316	150	2
GunPointMaleVersusFemale	135	316	150	2
GunPointAgeSpan	135	316	150	2
UMD	36	144	150	3
Epilepsy2	80	11420	178	2

Continued on next page

Dataset Name	Train Size	Test Size	Length	Classes
Colposcopy	100	100	180	6
Fungi	18	186	201	18
WalkingSittingStanding	7352	2947	206	6
Epilepsy	137	138	207	4
Strawberry	613	370	235	2
ElectricDeviceDetection	623	3767	256	2
FiftyWords	450	455	270	50
ToeSegmentation1	40	228	277	2
DodgerLoopWeekend	20	138	288	2
DodgerLoopGame	20	138	288	2
DodgerLoopDay	78	80	288	7
CricketZ	390	390	300	12
CricketY	390	390	300	12
CricketX	390	390	300	12
FreezerRegularTrain	150	2850	301	2
FreezerSmallTrain	28	2850	301	2
UWaveGestureLibraryZ	896	3582	315	8
UWaveGestureLibraryY	896	3582	315	8
UWaveGestureLibraryX	896	3582	315	8
UWaveGestureLibrary	2238	3582	315	8
ToeSegmentation2	36	130	343	2
GestureMidAirD3	208	130	360	26
GestureMidAirD2	208	130	360	26
GestureMidAirD1	208	130	360	26
Symbols	25	995	398	6
HandMovementDirection	160	74	400	4
Heartbeat	204	205	405	2
OSULeaf	200	242	427	6
Fish	175	175	463	7

Appendix B Pairwise Accuracy Comparisons

B.1 Randomized Prototype and Distance Metric Parameter Testing




Figure 17: Random Prototype Distance Metrics Comparison 1

B.1 Randomized Prototype and Distance Metric Parameter Testing




Figure 18: Random Prototype Distance Metrics Comparison 2




Figure 19: Random Prototype Distance Metrics Comparison 3

B.2 Stratified Random Prototype and Distance Metric Parameter Testing

B.2 Stratified Random Prototype and Distance Metric Parameter Testing




Figure 20: Stratified Random Prototype Distance Metrics Comparison 1

B.2 Stratified Random Prototype and Distance Metric Parameter Testing




Figure 21: Stratified Random Prototype Distance Metrics Comparison 2

B.2 Stratified Random Prototype and Distance Metric Parameter Testing




Figure 22: Stratified Random Prototype Distance Metrics Comparison 3

B.3 Simple Ensemble Comparisons

B.3 Simple Ensemble Comparisons




Figure 23: Distance Ensemble Comparison 1

B.3 Simple Ensemble Comparisons




Figure 24: Distance Ensemble Comparison 2

B.3 Simple Ensemble Comparisons




Figure 25: Distance Ensemble Comparison 3

B.3 Simple Ensemble Comparisons




Figure 26: Distance Ensemble Comparison 4

B.3 Simple Ensemble Comparisons




Figure 27: Distance Ensemble Comparison 5

B.3 Simple Ensemble Comparisons




Figure 28: Distance Ensemble Comparison 6

B.3 Simple Ensemble Comparisons




Figure 29: Distance Ensemble Comparison 7

B.3 Simple Ensemble Comparisons




Figure 30: Distance Ensemble Comparison 8




Figure 31: Distance Ensemble Comparison 9

B.4 MSM Scaling Comparisons

B.4 MSM Scaling Comparisons




Figure 32: MSM Scaling Comparison 1

B.4 MSM Scaling Comparisons




Figure 33: MSM Scaling Comparison 2

B.4 MSM Scaling Comparisons




Figure 34: MSM Scaling Comparison 3

B.4 MSM Scaling Comparisons




Figure 35: MSM Scaling Comparison 4

B.4 MSM Scaling Comparisons




Figure 36: MSM Scaling Comparison 5

B.5 Large Scale MSM Scaling Comparisons




Figure 37: MSM Large Scale Comparison 1

B.5 Large Scale MSM Scaling Comparisons




Figure 38: MSM Large Scale Comparison 2




Figure 39: MSM Large Scale Comparison 3

Appendix C Wall Clock Times for Large Scale Tests

C.1 MSM Large Scale Testing

Table 9: Average Wall Clock Time For All Algorithms on All Datasets

Dataset	sprocket	multirocket	hydra	quant	mr_hy	mr_sp	hy_sp	mr_hy_sp
DuckDuckGeese	2663.19	6.90	0.26	6.85	6.85	2666.50	2660.18	2666.75
PEMS-SF	3380.17	6.53	1.31	17.65	7.06	3383.70	3377.74	3383.36
MindReading	4263.92	7.28	1.78	1145.62	7.84	4267.04	4263.71	4269.41
PhonemeSpectra	1815.93	30.70	9.87	56.76	70.01	1858.89	1819.89	1869.49
ShapeletSim	59.93	0.74	0.55	0.39	1.05	60.43	60.36	60.92
EMOPain	635.83	4.37	3.05	14.35	5.67	637.79	636.16	638.96
SmoothSubspace	0.18	0.43	0.17	0.23	0.47	0.52	0.32	0.62
MelbournePedestrian	5.50	4.81	1.19	0.84	4.67	7.49	4.72	7.83
ItalyPowerDemand	0.71	0.78	0.17	0.35	0.89	1.45	0.86	1.60
Chinatown	0.26	0.33	0.14	0.21	0.40	0.53	0.37	0.66
JapaneseVowels	3.99	2.18	1.61	0.59	1.44	4.49	3.78	4.62
RacketSports	2.99	0.43	0.21	0.34	0.46	1.27	1.10	1.40
LSST	35.56	9.91	2.13	10.49	11.82	41.89	33.03	43.41
Libras	1.36	1.51	0.84	0.39	1.54	2.43	2.00	2.13
FingerMovements	17.34	1.45	1.41	2.06	1.35	15.71	14.93	15.73
NATOPS	9.79	1.88	1.10	0.89	1.73	10.38	9.70	10.22
SharePriceIncrease	9.05	3.61	2.09	1.91	3.26	10.10	8.26	10.79
SyntheticControl	2.89	1.54	0.99	0.51	2.40	3.97	2.91	3.78
SonyAIBORobotSurface2	2.59	0.83	0.29	0.46	1.04	3.36	2.86	3.65
ERing	1.92	0.38	0.17	0.85	0.48	2.23	2.06	2.39
SonyAIBORobotSurface1	1.82	0.55	0.23	0.43	0.70	2.30	2.04	2.53
PhalangesOutlinesCorrect	20.97	6.57	1.99	4.39	6.99	23.63	18.77	24.06
ProximalPhalanxOutlineCorrect	6.47	3.86	2.28	1.13	2.11	6.36	5.07	8.21
MiddlePhalanxOutlineCorrect	6.55	1.94	1.96	1.07	3.13	7.27	5.87	6.97

Dataset	sprocket	multirocket	hydra	quant	mr_hy	mr_sp	hy_sp	mr_hy_sp
DistalPhalanxOutlineCorrect	6.49	3.16	0.48	1.01	1.57	6.27	5.87	7.25
ProximalPhalanxTW	4.33	1.94	1.19	0.63	2.06	4.40	4.22	5.14
ProximalPhalanxOutlineAgeGroup	4.43	1.95	1.91	0.80	1.70	4.93	4.48	5.16
MiddlePhalanxOutlineAgeGroup	4.06	1.46	0.70	0.64	1.27	4.53	4.01	5.25
MiddlePhalanxTW	3.49	1.61	1.06	0.78	2.12	4.37	4.25	4.90
DistalPhalanxTW	3.95	1.54	1.47	0.63	2.21	4.79	3.64	4.33
DistalPhalanxOutlineAgeGroup	3.98	1.34	0.93	0.72	1.28	3.90	3.38	4.42
TwoLeadECG	4.32	1.04	0.37	0.32	1.34	5.29	4.68	5.66
MoteStrain	4.90	1.11	0.40	0.32	1.44	5.95	5.29	6.35
ECG200	1.44	0.50	0.33	0.31	0.61	1.70	1.67	1.88
MedicalImages	10.59	2.95	2.20	0.88	3.16	10.77	9.58	10.71
BasicMotions	3.10	0.14	0.13	0.31	0.21	1.51	1.54	1.64
TwoPatterns	68.07	7.10	3.40	1.89	9.29	73.67	69.42	75.47
CBF	7.39	1.02	0.43	0.35	1.35	8.34	7.81	8.76
SwedishLeaf	15.76	2.43	1.42	0.97	3.25	16.35	15.91	17.09
BME	1.49	0.28	0.43	0.23	0.64	1.71	1.90	2.13
EyesOpenShut	6.71	0.28	0.15	0.49	0.28	6.84	6.82	6.98
FacesUCR	25.37	2.97	1.99	0.60	4.22	27.20	26.31	28.54
FaceAll	32.09	3.83	1.99	1.35	5.01	34.00	32.33	34.90
ECGFiveDays	7.66	0.96	0.76	0.99	1.63	8.54	8.42	9.31
ECG5000	76.16	5.36	4.11	1.00	8.18	81.29	78.41	83.57
ArticularyWordRecognition	52.13	2.23	1.84	2.50	2.51	51.60	50.90	51.77
PowerCons	5.40	1.25	1.05	0.39	1.60	6.01	5.65	5.87
Plane	2.95	0.63	0.51	0.33	0.69	3.41	3.25	3.50
GunPointOldVersusYoung	7.11	1.48	1.61	0.31	1.78	7.71	7.32	7.63
GunPointMaleVersusFemale	6.66	1.72	1.53	0.38	2.48	7.61	7.73	8.17

Continued on next page

C.1 *MSM Large Scale Testing*

Dataset	sprocket	multirocket	hydra	quant	mr_hy	mr_sp	hy_sp	mr_hy_sp
GunPointAgeSpan	7.10	1.83	1.96	0.39	2.27	7.70	7.43	7.85
GunPoint	2.16	0.35	0.17	0.30	0.42	2.43	2.31	2.59
UMD	2.01	0.30	0.17	0.31	0.39	2.23	2.16	2.40
Wafer	128.71	9.43	7.29	1.17	17.15	140.25	131.39	144.28
Handwriting	24.27	1.59	0.70	1.04	2.03	25.60	24.84	26.20
ChlorineConcentration	92.67	6.12	3.72	1.90	8.52	95.98	93.75	98.69
Adiac	18.44	2.19	2.12	1.39	2.24	18.18	17.84	18.82
Epilepsy2	218.31	24.36	8.59	0.77	89.77	288.70	226.87	306.67
Colposcopy	3.90	0.72	0.38	0.40	0.75	4.29	4.04	4.46
Fungi	3.28	0.36	0.24	0.32	0.48	3.53	3.49	3.75
Epilepsy	14.63	1.70	1.18	0.59	1.76	13.59	13.12	13.77
Wine	2.82	0.25	0.16	0.33	0.29	2.97	2.95	3.12
Strawberry	39.42	3.05	2.64	1.37	3.23	39.75	38.56	40.49
ArrowHead	5.57	0.36	0.26	0.30	0.49	5.81	5.81	6.06
ElectricDeviceDetection	215.68	4.91	5.44	1.79	8.89	220.16	219.67	223.29
WordSynonyms	49.14	2.42	1.81	0.85	2.47	49.72	49.60	50.61
FiftyWords	47.88	2.19	1.59	1.82	2.58	48.67	48.63	50.48
Trace	8.83	0.64	0.42	0.37	0.77	9.21	9.12	9.50
ToeSegmentation1	9.72	0.44	0.39	0.28	0.69	10.05	10.08	10.42
DodgerLoopWeekend	5.60	0.35	0.26	0.26	0.46	5.81	5.83	6.06
DodgerLoopGame	6.20	0.33	0.25	0.23	0.45	6.40	6.42	6.64
DodgerLoopDay	7.85	0.35	0.22	0.31	0.44	8.08	8.03	8.28
CricketZ	55.73	2.53	2.01	1.14	2.97	56.07	55.42	56.75
CricketY	56.56	1.75	1.40	1.08	3.23	57.98	57.75	59.32
CricketX	56.02	2.07	1.58	1.15	2.91	57.22	56.68	57.93
FreezerRegularTrain	168.66	4.39	3.89	0.67	8.03	172.79	172.46	176.60

Continued on next page

Dataset	sprocket	multirocket	hydra	quant	mr_hy	mr_sp	hy_sp	mr_hy_sp
FreezerSmallTrain	134.02	4.01	3.58	0.55	7.46	137.89	137.60	141.47
UWaveGestureLibraryZ	335.95	6.22	5.91	2.90	12.75	342.19	341.39	347.15
UWaveGestureLibraryY	358.37	6.13	6.21	3.19	12.80	364.54	363.31	370.48
UWaveGestureLibraryX	361.95	6.40	5.98	2.97	13.94	369.29	368.64	375.56
UWaveGestureLibrary	47.53	1.47	1.30	0.69	2.04	48.76	48.41	49.11
Lightning7	9.76	0.37	0.23	0.35	0.46	10.00	9.95	10.22
ToeSegmentation2	9.64	0.37	0.29	0.29	0.52	9.87	9.91	10.16
DiatomSizeReduction	18.13	0.53	0.52	0.32	0.91	18.52	18.63	19.04
FaceFour	7.52	0.31	0.21	0.27	0.37	7.68	7.70	7.88
GestureMidAirD3	30.59	1.61	1.79	0.86	2.43	31.93	31.10	32.17
GestureMidAirD2	33.72	1.84	2.08	0.88	2.55	34.52	34.01	35.41
GestureMidAirD1	33.03	1.31	1.31	1.14	2.31	33.63	33.78	34.47
Symbols	116.93	1.52	1.76	0.42	3.14	118.32	118.68	120.07
HandMovementDirection	198.87	1.01	0.57	1.34	1.24	199.54	199.34	200.07
Heartbeat	2193.58	2.62	2.23	4.10	3.83	2194.98	2194.59	2195.90
Yoga	803.03	5.64	6.02	1.16	10.14	807.20	807.84	812.77
OSULeaf	81.53	1.63	1.96	0.62	2.77	82.77	82.67	83.50
Meat	21.25	0.32	0.23	0.32	0.40	21.44	21.45	21.66
Fish	92.63	1.40	1.72	0.62	2.44	93.63	93.33	94.54
FordA	2472.92	22.51	13.86	16.51	40.60	2494.93	2483.25	2508.32
FordB	2369.59	71.60	17.58	24.21	97.10	2421.89	2376.32	2436.19
Ham	45.46	0.78	1.01	0.64	1.26	46.01	46.03	46.66

C.2 Euclidean Large Scale Testing

[H]

Table 10: Average Wall Clock Time For All Algorithms on All Datasets with Euclidean SPROCKET

Dataset	sprocket	multirocket	hydra	quant	mr_hy	mr_sp	hy_sp	mr_hy_sp
DuckDuckGeese	7.29	7.78	0.28	7.06	7.62	11.03	3.89	11.29
PEMS-SF	10.63	7.74	1.03	15.36	8.04	14.19	8.26	15.44
MindReading	7.74	10.00	2.18	1085.49	10.61	13.10	7.74	15.19
PhonemeSpectra	14.08	56.10	10.38	62.79	92.14	76.54	16.06	83.04
ShapeletSim	0.11	0.80	0.56	0.47	1.11	0.71	0.58	1.20
EMOPain	2.73	4.40	2.39	13.73	5.86	4.80	3.70	6.76
SmoothSubspace	0.10	0.46	0.18	0.23	0.45	0.41	0.24	0.51
MelbournePedestrian	2.05	4.50	1.06	0.92	5.46	5.04	1.16	6.35
ItalyPowerDemand	0.12	0.81	0.15	0.34	0.86	0.84	0.24	0.99
Chinatown	0.08	0.31	0.13	0.21	0.38	0.35	0.18	0.45
JapaneseVowels	0.79	1.48	1.22	1.11	1.77	1.97	0.77	1.83
RacketSports	2.47	0.45	0.20	0.39	0.49	0.42	0.25	0.56
LSST	6.25	13.58	3.20	12.92	19.47	18.17	3.28	19.92
Libras	0.62	1.45	1.13	0.44	1.58	1.52	1.13	1.32
FingerMovements	2.38	1.30	1.97	1.61	1.84	2.10	0.97	1.67
NATOPS	1.04	1.52	1.12	0.84	1.56	1.65	1.22	1.70
SharePriceIncrease	1.56	3.16	1.36	1.80	4.30	3.66	2.48	4.60
SyntheticControl	0.80	2.24	0.93	0.43	2.18	1.40	0.70	1.77
SonyAIBORobotSurface2	0.15	0.89	0.33	0.48	1.12	0.95	0.45	1.26
ERing	0.12	0.35	0.20	0.53	0.44	0.36	0.28	0.56
SonyAIBORobotSurface1	0.10	0.58	0.24	0.27	0.74	0.61	0.32	0.85

Continued on next page

Dataset	sprocket	multirocket	hydra	quant	mr_hy	mr_sp	hy_sp	mr_hy_sp
PhalangesOutlinesCorrect	3.21	7.06	2.35	3.75	8.81	10.77	4.96	9.37
ProximalPhalanxOutlineCorrect	1.55	3.02	2.56	1.12	2.79	2.14	1.64	2.84
MiddlePhalanxOutlineCorrect	2.36	2.99	1.22	1.24	1.90	2.30	1.01	3.39
DistalPhalanxOutlineCorrect	1.64	2.77	1.21	1.07	3.02	2.02	0.67	2.57
ProximalPhalanxTW	1.12	1.65	1.60	0.58	2.61	2.05	0.63	1.74
ProximalPhalanxOutlineAgeGroup	1.10	2.05	1.45	0.70	2.10	2.17	0.70	1.51
MiddlePhalanxOutlineAgeGroup	1.56	2.20	0.87	0.68	1.15	2.44	0.88	2.13
MiddlePhalanxTW	1.07	2.72	1.14	0.76	1.59	1.67	0.82	1.91
DistalPhalanxTW	1.32	2.35	0.95	0.64	1.76	1.04	0.48	1.32
DistalPhalanxOutlineAgeGroup	0.99	1.30	1.13	0.62	1.66	1.53	0.73	1.45
TwoLeadECG	0.15	1.10	0.39	0.32	1.41	1.18	0.52	1.55
MoteStrain	0.14	1.16	0.42	0.37	1.50	1.24	0.53	1.64
ECG200	0.21	0.51	0.33	0.36	0.53	0.49	0.40	0.63
MedicalImages	1.37	2.96	1.43	0.92	2.01	1.89	1.81	2.72
BasicMotions	1.87	0.17	0.14	0.35	0.22	0.12	0.15	0.26
TwoPatterns	1.64	8.72	2.66	2.13	15.27	13.58	3.32	15.71
CBF	0.11	1.09	0.44	0.32	1.43	1.10	0.54	1.53
SwedishLeaf	1.37	2.85	2.03	0.97	3.35	2.49	1.01	2.24
BME	0.07	0.31	0.18	0.28	0.39	0.29	0.22	0.46
EyesOpenShut	0.09	0.30	0.15	0.51	0.29	0.23	0.20	0.37
FacesUCR	1.01	2.98	2.42	0.54	4.52	3.44	2.57	4.54
FaceAll	1.93	3.57	2.12	2.11	5.53	4.49	3.15	5.91
ECGFiveDays	0.14	0.91	0.52	0.31	1.35	0.96	0.64	1.49
ECG5000	1.68	7.12	3.89	1.07	12.85	10.46	3.61	12.74
ArticularyWordRecognition	2.41	1.66	1.54	2.92	1.98	1.76	1.37	2.32
PowerCons	0.32	1.34	1.29	0.37	1.94	1.23	0.93	1.55

Continued on next page

Dataset	sprocket	multirocket	hydra	quant	mr_hy	mr_sp	hy_sp	mr_hy_sp
Plane	0.25	0.59	0.41	0.31	0.65	0.57	0.46	0.78
GunPointOldVersusYoung	0.87	1.55	1.74	0.29	2.23	2.07	1.63	1.78
GunPointMaleVersusFemale	0.91	1.90	1.64	0.38	2.10	1.62	1.46	2.05
GunPointAgeSpan	1.05	2.14	1.69	0.41	2.47	1.96	1.38	2.09
GunPoint	0.11	0.33	0.18	0.30	0.39	0.35	0.25	0.51
UMD	0.07	0.32	0.18	0.30	0.38	0.29	0.22	0.45
Wafer	3.01	15.77	4.03	1.21	22.83	18.19	5.18	23.60
Handwriting	0.48	1.73	0.80	1.24	2.19	1.91	1.09	2.55
ChlorineConcentration	1.80	6.42	3.12	2.18	11.03	8.43	3.76	11.07
Adiac	1.47	2.31	1.69	1.54	2.97	2.62	1.91	2.57
Epilepsy2	1.71	50.86	21.61	0.96	174.55	131.00	23.19	172.58
Colposcopy	0.20	0.94	0.42	0.59	0.95	0.76	0.45	1.02
Fungi	0.05	0.40	0.25	0.38	0.52	0.33	0.27	0.58
Epilepsy	3.48	1.69	1.01	0.62	1.77	1.62	1.12	1.93
Wine	0.06	0.30	0.16	0.35	0.33	0.23	0.19	0.40
Strawberry	1.71	2.84	3.37	1.30	3.14	2.95	2.53	3.70
ArrowHead	0.09	0.42	0.26	0.28	0.54	0.38	0.33	0.64
ElectricDeviceDetection	1.55	6.05	4.41	1.09	11.67	8.14	5.67	12.04
WordSynonyms	1.00	2.73	1.69	0.85	2.68	2.07	2.02	2.56
FiftyWords	1.73	2.47	1.83	1.64	2.60	2.09	1.73	3.23
Trace	0.23	0.72	0.52	0.38	0.84	0.62	0.61	0.96
ToeSegmentation1	0.12	0.51	0.38	0.39	0.79	0.53	0.47	0.91
DodgerLoopWeekend	0.09	0.39	0.26	0.31	0.49	0.31	0.30	0.55
DodgerLoopGame	0.08	0.39	0.27	0.32	0.50	0.30	0.31	0.56
DodgerLoopDay	0.09	0.41	0.23	0.37	0.48	0.33	0.27	0.56
CricketZ	2.14	2.80	2.11	1.25	2.82	2.65	2.08	3.94

Continued on next page

Dataset	sprocket	multirocket	hydra	quant	mr_hy	mr_sp	hy_sp	mr_hy_sp
CricketY	0.88	1.90	1.64	1.23	3.04	1.82	1.56	3.16
CricketX	1.49	2.65	1.62	1.17	3.23	2.53	2.04	2.94
FreezerRegularTrain	0.84	4.84	3.86	0.74	9.01	5.84	4.66	9.82
FreezerSmallTrain	0.71	4.81	3.67	0.64	9.61	6.12	4.37	10.03
UWaveGestureLibraryZ	3.38	9.10	6.09	2.93	18.16	13.66	8.46	19.47
UWaveGestureLibraryY	2.38	9.74	6.68	3.07	19.15	14.08	11.86	19.48
UWaveGestureLibraryX	2.20	8.71	7.03	2.54	16.52	12.64	7.35	18.76
UWaveGestureLibrary	0.70	1.57	1.16	0.61	2.03	1.53	1.33	1.92
Lightning7	0.07	0.36	0.23	0.35	0.47	0.33	0.27	0.55
ToeSegmentation2	0.08	0.37	0.29	0.33	0.52	0.31	0.34	0.60
DiatomSizeReduction	0.12	0.52	0.51	0.27	0.88	0.49	0.61	0.99
FaceFour	0.07	0.31	0.21	0.28	0.38	0.23	0.25	0.44
GestureMidAirD3	0.51	1.67	1.53	0.87	1.89	1.26	1.62	2.31
GestureMidAirD2	1.01	1.82	1.86	1.61	2.40	1.59	1.46	2.41
GestureMidAirD1	0.95	1.77	1.77	1.77	2.40	1.75	1.54	2.07
Symbols	0.35	1.65	1.77	0.47	3.26	1.85	2.11	3.61
HandMovementDirection	0.49	1.10	0.62	1.71	1.31	1.18	0.97	1.71
Heartbeat	2.40	3.13	2.12	4.45	3.27	3.39	2.72	4.53
Yoga	2.79	5.69	6.92	1.35	13.27	7.69	7.56	13.35
OSULeaf	0.97	1.86	1.85	0.69	2.89	2.18	1.51	2.46
Meat	0.07	0.35	0.23	0.33	0.41	0.27	0.27	0.49
Fish	0.74	1.61	1.95	0.63	2.78	1.75	1.56	2.19
FordA	12.01	35.30	16.22	18.09	62.56	49.73	18.56	72.50
FordB	12.82	43.55	22.24	26.76	94.80	109.70	23.77	107.19
Ham	0.29	1.02	1.04	0.58	1.66	1.08	1.00	1.57